Observation-wise versus event-level scoring
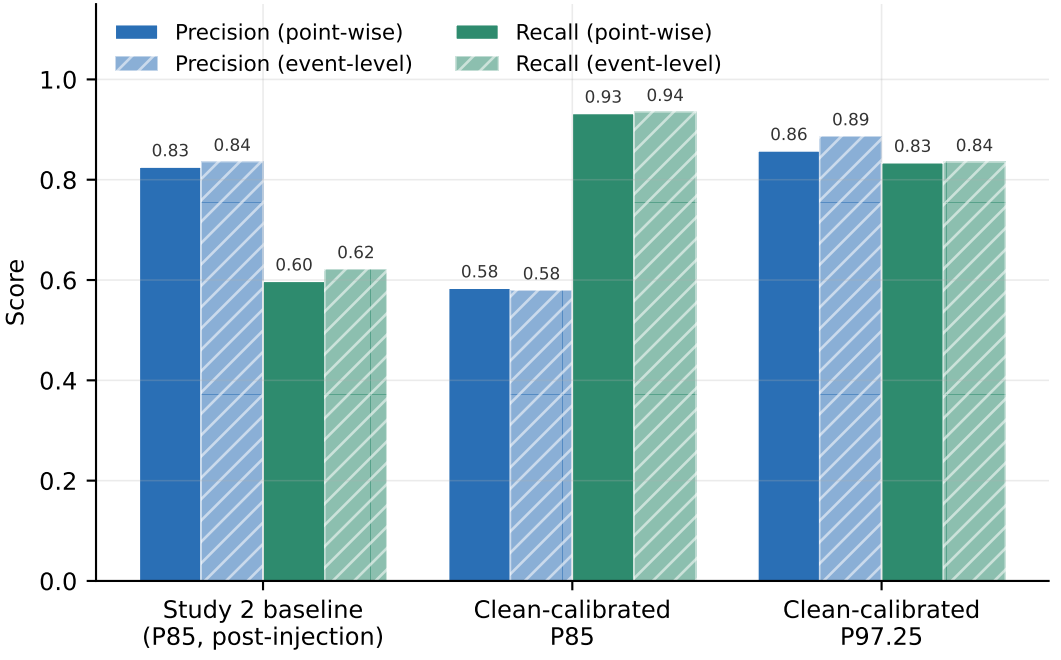
Identical predictions, two metrics. Event-level scoring (Tatbul et al., α=0.5) asks whether the event was caught at all rather than whether every observation of it was flagged. The two nearly coincide here: persistence requires 2 consecutive exceedances and events last 2 observations, so detection is all-or-nothing and there is almost no partial credit to award.